

01_eda

April 21, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import sys
from pathlib import Path

# Add src directory to sys.path to import config
# This assumes the notebook is in RECSYS_FINAL/notebooks/
project_root = Path.cwd().parent # Should be RECSYS_FINAL
src_path = project_root / "src"
sys.path.append(str(src_path))

# Import config variables
import config

# Set some display options for pandas
pd.set_option('display.max_columns', 50)
pd.set_option('display.max_rows', 100)

# Plotting style
sns.set_style("whitegrid")
plt.rcParams['figure.figsize'] = (12, 6)

print(f"Project Root: {project_root}")
print(f"Configured Raw Data Path: {config.RAW_DATA_DIR}")
print(f"All Raw Files Found: {config.check_raw_data_exists()}")
```

```
Loading .env from: /Users/mohit/Desktop/everything/ATLAS/Semester
4/Pinnacle/recsys_final/.env
Database URI configured: Yes
Project Root: /Users/mohit/Desktop/everything/ATLAS/Semester
4/Pinnacle/recsys_final
Configured Raw Data Path: /Users/mohit/Desktop/everything/ATLAS/Semester
4/Pinnacle/recsys_final/data/raw
All raw data files found.
All Raw Files Found: True
```

```
[2]: # Load all datasets using paths from config
try:
    assessments_df = pd.read_csv(config.ASSESSMENTS_CSV)
    courses_df = pd.read_csv(config.COURSES_CSV)
    student_assessment_df = pd.read_csv(config.STUDENT_ASSESSMENT_CSV)
    student_info_df = pd.read_csv(config.STUDENT_INFO_CSV)
    student_registration_df = pd.read_csv(config.STUDENT_REGISTRATION_CSV)
    student_vle_df = pd.read_csv(config.STUDENT_VLE_CSV)
    vle_df = pd.read_csv(config.VLE_CSV)
    print("All CSV files loaded successfully.")
except FileNotFoundError as e:
    print(f"Error loading files: {e}")
    print("Please ensure the CSV files are in the data/raw/ directory.")
    # Stop execution or handle appropriately

# Store dataframes in a dictionary for easier access
dataframes = {
    "assessments": assessments_df,
    "courses": courses_df,
    "student_assessment": student_assessment_df,
    "student_info": student_info_df,
    "student_registration": student_registration_df,
    "student_vle": student_vle_df,
    "vle": vle_df,
}
```

All CSV files loaded successfully.

```
[3]: # Basic inspection of each dataframe
for name, df in dataframes.items():
    print(f"--- DataFrame: {name} ---")
    print(f"Shape: {df.shape}")
    print("Info:")
    df.info()
    print("\nHead:")
    print(df.head())
    print("\nMissing Values:")
    print(df.isnull().sum())
    # Only show describe() for dataframes with numerical columns
    numeric_cols = df.select_dtypes(include=np.number).columns
    if len(numeric_cols) > 0:
        print("\nDescribe (Numerical):")
        print(df.describe())
    # Describe categorical columns
    categorical_cols = df.select_dtypes(include=['object', 'category']).columns
    if len(categorical_cols) > 0:
        print("\nDescribe (Categorical):")
        print(df.describe(include=['object', 'category']))
```

```
print("-" * (len(name) + 22))
print("\n")
```

--- DataFrame: assessments ---

Shape: (206, 6)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 206 entries, 0 to 205

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	code_module	206 non-null	object
1	code_presentation	206 non-null	object
2	id_assessment	206 non-null	int64
3	assessment_type	206 non-null	object
4	date	195 non-null	float64
5	weight	206 non-null	float64

dtypes: float64(2), int64(1), object(3)

memory usage: 9.8+ KB

Head:

	code_module	code_presentation	id_assessment	assessment_type	date	weight
0	AAA	2013J	1752	TMA	19.0	10.0
1	AAA	2013J	1753	TMA	54.0	20.0
2	AAA	2013J	1754	TMA	117.0	20.0
3	AAA	2013J	1755	TMA	166.0	20.0
4	AAA	2013J	1756	TMA	215.0	30.0

Missing Values:

code_module	0
code_presentation	0
id_assessment	0
assessment_type	0
date	11
weight	0

dtype: int64

Describe (Numerical):

	id_assessment	date	weight
count	206.000000	195.000000	206.000000
mean	26473.975728	145.005128	20.873786
std	10098.625521	76.001119	30.384224
min	1752.000000	12.000000	0.000000
25%	15023.250000	71.000000	0.000000
50%	25364.500000	152.000000	12.500000
75%	34891.750000	222.000000	24.250000
max	40088.000000	261.000000	100.000000

Describe (Categorical):

	code_module	code_presentation	assessment_type
count	206	206	206
unique	7	4	3
top	FFF	2014J	TMA
freq	52	57	106

--- DataFrame: courses ---

Shape: (22, 3)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 22 entries, 0 to 21

Data columns (total 3 columns):

#	Column	Non-Null Count	Dtype
0	code_module	22 non-null	object
1	code_presentation	22 non-null	object
2	module_presentation_length	22 non-null	int64

dtypes: int64(1), object(2)

memory usage: 660.0+ bytes

Head:

	code_module	code_presentation	module_presentation_length
0	AAA	2013J	268
1	AAA	2014J	269
2	BBB	2013J	268
3	BBB	2014J	262
4	BBB	2013B	240

Missing Values:

code_module 0

code_presentation 0

module_presentation_length 0

dtype: int64

Describe (Numerical):

	module_presentation_length
count	22.000000
mean	255.545455
std	13.654677
min	234.000000
25%	241.000000
50%	261.500000
75%	268.000000
max	269.000000

Describe (Categorical):

	code_module	code_presentation
count	22	22
unique	7	4
top	BBB	2014J
freq	4	7

--- DataFrame: student_assessment ---

Shape: (173912, 5)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 173912 entries, 0 to 173911

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	id_assessment	173912 non-null	int64
1	id_student	173912 non-null	int64
2	date_submitted	173912 non-null	int64
3	is_banked	173912 non-null	int64
4	score	173739 non-null	float64

dtypes: float64(1), int64(4)

memory usage: 6.6 MB

Head:

	id_assessment	id_student	date_submitted	is_banked	score
0	1752	11391	18	0	78.0
1	1752	28400	22	0	70.0
2	1752	31604	17	0	72.0
3	1752	32885	26	0	69.0
4	1752	38053	19	0	79.0

Missing Values:

id_assessment	0
id_student	0
date_submitted	0
is_banked	0
score	173

dtype: int64

Describe (Numerical):

	id_assessment	id_student	date_submitted	is_banked	\
count	173912.000000	1.739120e+05	173912.000000	173912.000000	
mean	26553.803556	7.051507e+05	116.032942	0.010977	
std	8829.784254	5.523952e+05	71.484148	0.104194	
min	1752.000000	6.516000e+03	-11.000000	0.000000	
25%	15022.000000	5.044290e+05	51.000000	0.000000	

50%	25359.000000	5.852080e+05	116.000000	0.000000
75%	34883.000000	6.344980e+05	173.000000	0.000000
max	37443.000000	2.698588e+06	608.000000	1.000000

```

                score
count  173739.000000
mean      75.799573
std       18.798107
min        0.000000
25%       65.000000
50%       80.000000
75%       90.000000
max      100.000000

```

--- DataFrame: student_info ---

Shape: (32593, 12)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 32593 entries, 0 to 32592

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	code_module	32593 non-null	object
1	code_presentation	32593 non-null	object
2	id_student	32593 non-null	int64
3	gender	32593 non-null	object
4	region	32593 non-null	object
5	highest_education	32593 non-null	object
6	imd_band	31482 non-null	object
7	age_band	32593 non-null	object
8	num_of_prev_attempts	32593 non-null	int64
9	studied_credits	32593 non-null	int64
10	disability	32593 non-null	object
11	final_result	32593 non-null	object

dtypes: int64(3), object(9)

memory usage: 3.0+ MB

Head:

	code_module	code_presentation	id_student	gender	region \
0	AAA	2013J	11391	M	East Anglian Region
1	AAA	2013J	28400	F	Scotland
2	AAA	2013J	30268	F	North Western Region
3	AAA	2013J	31604	F	South East Region
4	AAA	2013J	32885	F	West Midlands Region

highest_education	imd_band	age_band	num_of_prev_attempts \
-------------------	----------	----------	------------------------

0	HE Qualification	90-100%	55<=	0
1	HE Qualification	20-30%	35-55	0
2	A Level or Equivalent	30-40%	35-55	0
3	A Level or Equivalent	50-60%	35-55	0
4	Lower Than A Level	50-60%	0-35	0

	studied_credits	disability	final_result
0	240	N	Pass
1	60	N	Pass
2	60	Y	Withdrawn
3	60	N	Pass
4	60	N	Pass

Missing Values:

code_module	0
code_presentation	0
id_student	0
gender	0
region	0
highest_education	0
imd_band	1111
age_band	0
num_of_prev_attempts	0
studied_credits	0
disability	0
final_result	0
dtype:	int64

Describe (Numerical):

	id_student	num_of_prev_attempts	studied_credits
count	3.259300e+04	32593.000000	32593.000000
mean	7.066877e+05	0.163225	79.758691
std	5.491673e+05	0.479758	41.071900
min	3.733000e+03	0.000000	30.000000
25%	5.085730e+05	0.000000	60.000000
50%	5.903100e+05	0.000000	60.000000
75%	6.444530e+05	0.000000	120.000000
max	2.716795e+06	6.000000	655.000000

Describe (Categorical):

	code_module	code_presentation	gender	region	highest_education \
count	32593	32593	32593	32593	32593
unique	7	4	2	13	5
top	BBB	2014J	M	Scotland	A Level or Equivalent
freq	7909	11260	17875	3446	14045

	imd_band	age_band	disability	final_result
count	31482	32593	32593	32593

unique	10	3	2	4
top	20-30%	0-35	N	Pass
freq	3654	22944	29429	12361

--- DataFrame: student_registration ---

Shape: (32593, 5)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 32593 entries, 0 to 32592

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	code_module	32593 non-null	object
1	code_presentation	32593 non-null	object
2	id_student	32593 non-null	int64
3	date_registration	32548 non-null	float64
4	date_unregistration	10072 non-null	float64

dtypes: float64(2), int64(1), object(2)

memory usage: 1.2+ MB

Head:

	code_module	code_presentation	id_student	date_registration	\
0	AAA	2013J	11391	-159.0	
1	AAA	2013J	28400	-53.0	
2	AAA	2013J	30268	-92.0	
3	AAA	2013J	31604	-52.0	
4	AAA	2013J	32885	-176.0	

	date_unregistration
0	NaN
1	NaN
2	12.0
3	NaN
4	NaN

Missing Values:

code_module	0
code_presentation	0
id_student	0
date_registration	45
date_unregistration	22521

dtype: int64

Describe (Numerical):

	id_student	date_registration	date_unregistration
count	3.259300e+04	32548.000000	10072.000000

mean	7.066877e+05	-69.411300	49.757645
std	5.491673e+05	49.260522	82.460890
min	3.733000e+03	-322.000000	-365.000000
25%	5.085730e+05	-100.000000	-2.000000
50%	5.903100e+05	-57.000000	27.000000
75%	6.444530e+05	-29.000000	109.000000
max	2.716795e+06	167.000000	444.000000

Describe (Categorical):

	code_module	code_presentation
count	32593	32593
unique	7	4
top	BBB	2014J
freq	7909	11260

--- DataFrame: student_vle ---

Shape: (10655280, 6)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10655280 entries, 0 to 10655279

Data columns (total 6 columns):

#	Column	Dtype
0	code_module	object
1	code_presentation	object
2	id_student	int64
3	id_site	int64
4	date	int64
5	sum_click	int64

dtypes: int64(4), object(2)

memory usage: 487.8+ MB

Head:

	code_module	code_presentation	id_student	id_site	date	sum_click
0	AAA	2013J	28400	546652	-10	4
1	AAA	2013J	28400	546652	-10	1
2	AAA	2013J	28400	546652	-10	1
3	AAA	2013J	28400	546614	-10	11
4	AAA	2013J	28400	546714	-10	1

Missing Values:

code_module	0
code_presentation	0
id_student	0
id_site	0
date	0

```
sum_click          0
dtype: int64
```

Describe (Numerical):

	id_student	id_site	date	sum_click
count	1.065528e+07	1.065528e+07	1.065528e+07	1.065528e+07
mean	7.333336e+05	7.383234e+05	9.517400e+01	3.716946e+00
std	5.827060e+05	1.312196e+05	7.607130e+01	8.849047e+00
min	6.516000e+03	5.267210e+05	-2.500000e+01	1.000000e+00
25%	5.077430e+05	6.735190e+05	2.500000e+01	1.000000e+00
50%	5.882360e+05	7.300690e+05	8.600000e+01	2.000000e+00
75%	6.464840e+05	8.770300e+05	1.560000e+02	3.000000e+00
max	2.698588e+06	1.049562e+06	2.690000e+02	6.977000e+03

Describe (Categorical):

	code_module	code_presentation
count	10655280	10655280
unique	7	4
top	FFF	2014J
freq	4014499	3619452

--- DataFrame: vle ---

Shape: (6364, 6)

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 6364 entries, 0 to 6363

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	id_site	6364 non-null	int64
1	code_module	6364 non-null	object
2	code_presentation	6364 non-null	object
3	activity_type	6364 non-null	object
4	week_from	1121 non-null	float64
5	week_to	1121 non-null	float64

dtypes: float64(2), int64(1), object(3)

memory usage: 298.4+ KB

Head:

	id_site	code_module	code_presentation	activity_type	week_from	week_to
0	546943	AAA	2013J	resource	NaN	NaN
1	546712	AAA	2013J	oucontent	NaN	NaN
2	546998	AAA	2013J	resource	NaN	NaN
3	546888	AAA	2013J	url	NaN	NaN
4	547035	AAA	2013J	resource	NaN	NaN

Missing Values:

id_site	0
code_module	0
code_presentation	0
activity_type	0
week_from	5243
week_to	5243

dtype: int64

Describe (Numerical):

	id_site	week_from	week_to
count	6.364000e+03	1121.000000	1121.000000
mean	7.260991e+05	15.204282	15.214987
std	1.283151e+05	8.792865	8.779806
min	5.267210e+05	0.000000	0.000000
25%	6.615928e+05	8.000000	8.000000
50%	7.300965e+05	15.000000	15.000000
75%	8.140162e+05	22.000000	22.000000
max	1.077905e+06	29.000000	29.000000

Describe (Categorical):

	code_module	code_presentation	activity_type
count	6364	6364	6364
unique	7	4	20
top	FFF	2013J	resource
freq	1967	1772	2660

```
[4]: print("--- Analyzing student_info ---")
student_info = dataframes['student_info']

# Distributions of categorical features
fig, axes = plt.subplots(3, 2, figsize=(15, 18))
sns.countplot(ax=axes[0, 0], x='gender', data=student_info)
axes[0, 0].set_title('Gender Distribution')

sns.countplot(ax=axes[0, 1], y='region', data=student_info,
              order=student_info['region'].value_counts().index)
axes[0, 1].set_title('Region Distribution')

sns.countplot(ax=axes[1, 0], y='highest_education', data=student_info,
              order=student_info['highest_education'].value_counts().index)
axes[1, 0].set_title('Highest Education Distribution')
```

```

sns.countplot(ax=axes[1, 1], y='imd_band', data=student_info,
    ↳order=student_info['imd_band'].value_counts(dropna=False).index) # Show NaNs
axes[1, 1].set_title('IMD Band Distribution (incl. NaN)')

sns.countplot(ax=axes[2, 0], x='age_band', data=student_info,
    ↳order=student_info['age_band'].value_counts().index)
axes[2, 0].set_title('Age Band Distribution')

sns.countplot(ax=axes[2, 1], x='disability', data=student_info)
axes[2, 1].set_title('Disability Status Distribution')

plt.tight_layout()
plt.show()


# Explore numerical features
print("\nPrevious Attempts Distribution:")
sns.histplot(student_info['num_of_prev_attempts'],
    ↳bins=range(student_info['num_of_prev_attempts'].max() + 2), kde=False)
plt.title('Distribution of Previous Attempts')
plt.xlabel('Number of Previous Attempts')
plt.ylabel('Count')
plt.show()

print("\nStudied Credits Distribution:")
sns.histplot(student_info['studied_credits'], bins=20, kde=True)
plt.title('Distribution of Studied Credits')
plt.xlabel('Studied Credits')
plt.ylabel('Count')
plt.show()

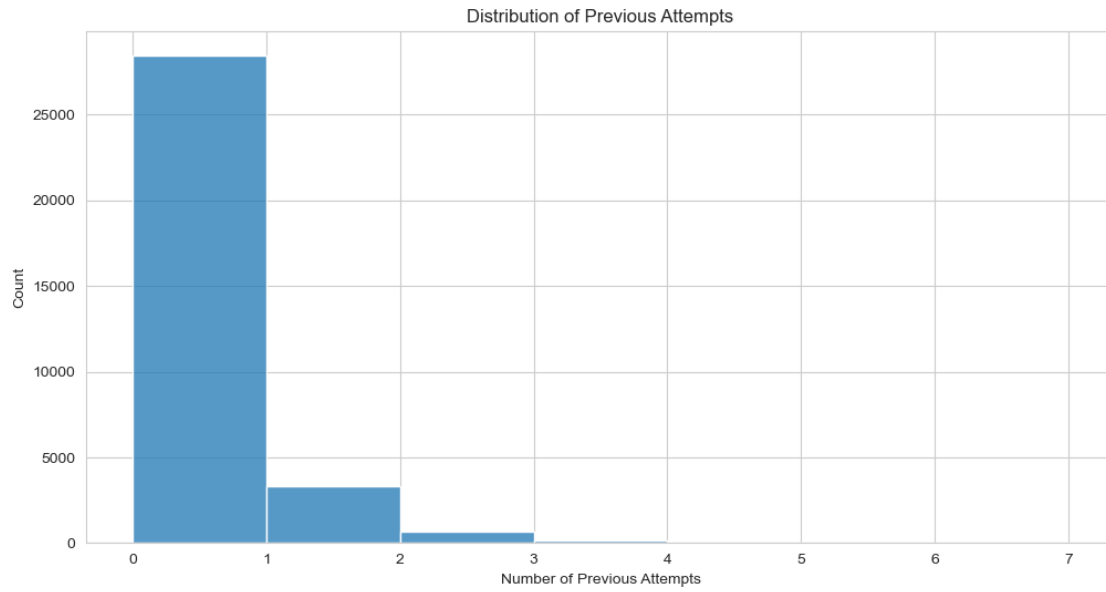
# Check final result distribution
print("\nFinal Result Distribution:")
sns.countplot(y='final_result', data=student_info,
    ↳order=student_info['final_result'].value_counts().index)
plt.title('Final Result Distribution')
plt.show()

```

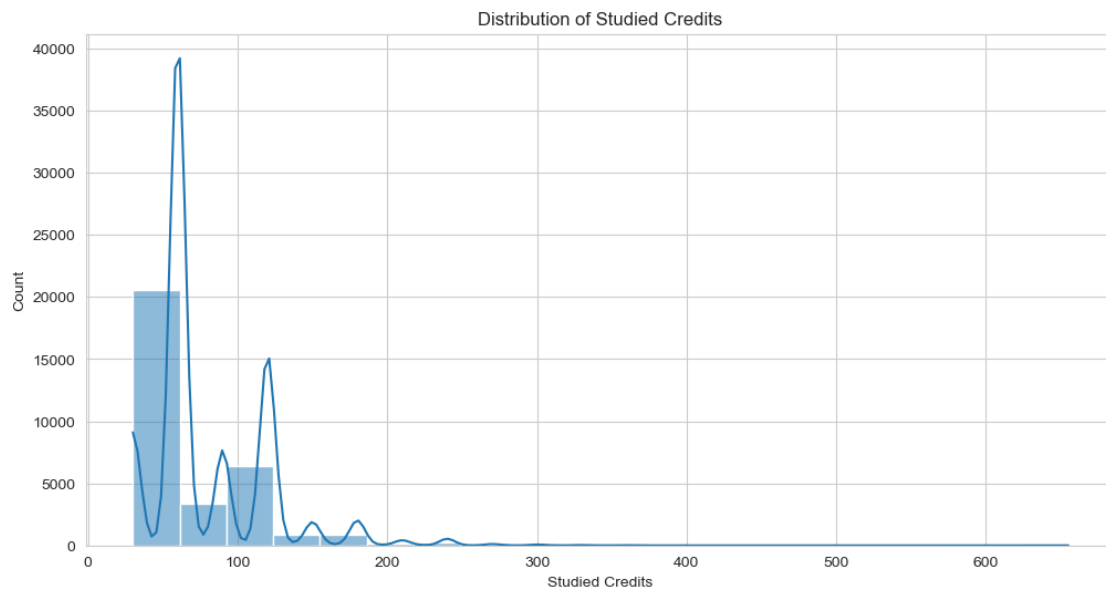
--- Analyzing student_info ---



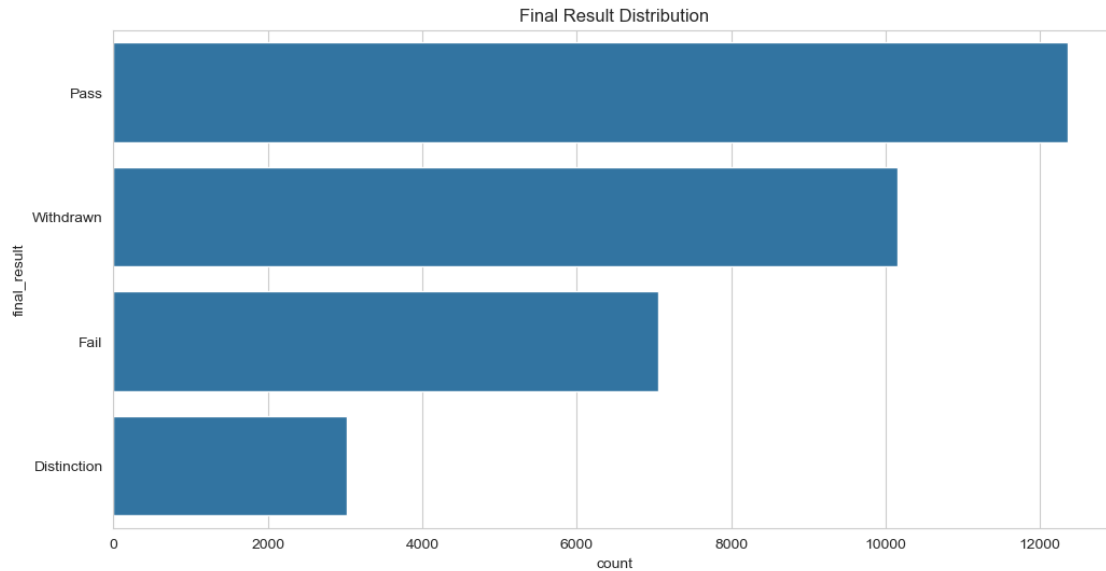
Previous Attempts Distribution:



Studied Credits Distribution:



Final Result Distribution:



```
[5]: print("--- Analyzing courses ---")
courses = dataframes['courses']

print(f"Number of unique modules: {courses['code_module'].nunique()}")
print(f"Number of unique presentations: {courses.shape[0]}") # Each row is a
↳ presentation
print("\nModules and their number of presentations:")
print(courses['code_module'].value_counts())

print("\nDistribution of Module Presentation Lengths:")
sns.histplot(courses['module_presentation_length'], bins=20, kde=False)
plt.title('Distribution of Module Presentation Lengths')
plt.xlabel('Length (days)')
plt.ylabel('Count')
plt.show()

# Create presentation_id for later use (might do this again in preprocessing)
courses['presentation_id'] = courses['code_module'] + '_' +
↳ courses['code_presentation']
print(f"\nCheck unique presentation IDs: {courses['presentation_id'].nunique()}")
↳ (should match shape[0])")
```

--- Analyzing courses ---

Number of unique modules: 7

Number of unique presentations: 22

Modules and their number of presentations:

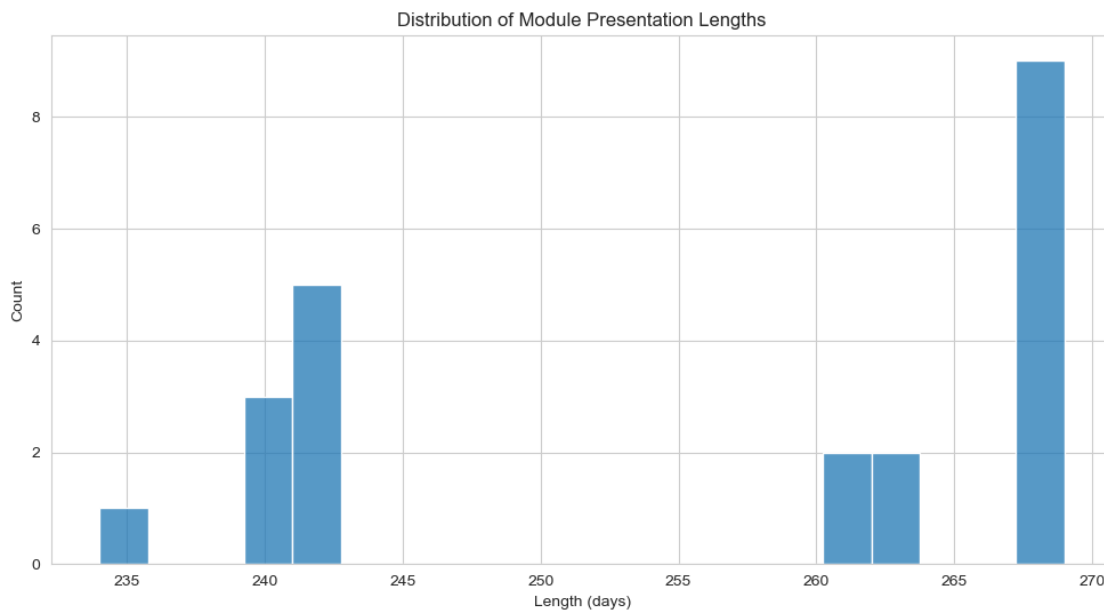
code_module

```

BBB    4
DDD    4
FFF    4
EEE    3
GGG    3
AAA    2
CCC    2
Name: count, dtype: int64

```

Distribution of Module Presentation Lengths:



Check unique presentation IDs: 22 (should match shape[0])

```

[6]: print("--- Analyzing student_registration ---")
      student_reg = dataframes['student_registration']

      # Check for multiple registrations per student per presentation (should be
      # unique ideally)
      reg_duplicates = student_reg.duplicated(subset=['id_student', 'code_module',
      'code_presentation']).sum()
      print(f"\nDuplicate registrations (same student, same presentation):
      {reg_duplicates}")

      # Distribution of registration dates (relative to presentation start, which is
      # 0)
      print("\nDistribution of Registration Dates:")

```



```

sns.histplot(student_reg['date_registration'].dropna(), bins=50, kde=True)
plt.title('Distribution of Registration Date (relative to start)')
plt.xlabel('Days Relative to Presentation Start')
plt.ylabel('Count')
plt.show()
# Note: Negative values mean registered before start, positive after.

# Distribution of unregistration dates
print("\nDistribution of Unregistration Dates:")
sns.histplot(student_reg['date_unregistration'].dropna(), bins=50, kde=True)
plt.title('Distribution of Unregistration Date (relative to start)')
plt.xlabel('Days Relative to Presentation Start')
plt.ylabel('Count')
plt.show()


print(f"\nPercentage of registrations with an unregistration date: \
{student_reg['date_unregistration'].notnull().mean() * 100:.2f}%")

```


--- Analyzing student_registration ---

Duplicate registrations (same student, same presentation): 0

Distribution of Registration Dates:



Distribution of Unregistration Dates:



Percentage of registrations with an unregistration date: 30.90%

```
[7]: print("--- Analyzing student_vle ---")
student_vle = dataframes['student_vle']

print(f"Total interaction records: {student_vle.shape[0]}")
print(f"Unique students in VLE logs: {student_vle['id_student'].nunique()}")
print(f"Unique VLE items interacted with: {student_vle['id_site'].nunique()}")

# Distribution of interaction dates
print("\nDistribution of Interaction Dates (relative to start):")
sns.histplot(student_vle['date'], bins=50, kde=False)
plt.title('Distribution of Interaction Dates')
plt.xlabel('Days Relative to Presentation Start')
plt.ylabel('Number of Interactions')
plt.show()

# Distribution of sum_click
# Handle potential outliers by looking at quantiles
print("\nDistribution of Clicks per Interaction Record:")
print(student_vle['sum_click'].describe(percentiles=[.25, .5, .75, .9, .95, .
↪99]))

# Plotting might be skewed, consider log scale or capping
sns.histplot(student_vle['sum_click'], bins=50, log_scale=(False, True)) # Log_
↪scale for y-axis
plt.title('Distribution of Clicks per Interaction Record (Log Y scale)')
```

```

plt.xlabel('Number of Clicks (sum_click)')
plt.ylabel('Frequency (Log Scale)')
plt.show()

# Interactions per student (across all their registered courses)
interactions_per_student = student_vle.groupby('id_student').size()
print("\nInteractions per Student (Summary Stats):")
print(interactions_per_student.describe())
sns.histplot(interactions_per_student, bins=50, log_scale=True)
plt.title('Distribution of Total Interaction Records per Student (Log Scale)')
plt.xlabel('Number of Interaction Records (Log Scale)')
plt.ylabel('Number of Students (Log Scale)')
plt.show()

# Clicks per student
clicks_per_student = student_vle.groupby('id_student')['sum_click'].sum()
print("\nTotal Clicks per Student (Summary Stats):")
print(clicks_per_student.describe())
sns.histplot(clicks_per_student, bins=50, log_scale=True)
plt.title('Distribution of Total Clicks per Student (Log Scale)')
plt.xlabel('Total Clicks (Log Scale)')
plt.ylabel('Number of Students (Log Scale)')
plt.show()

```

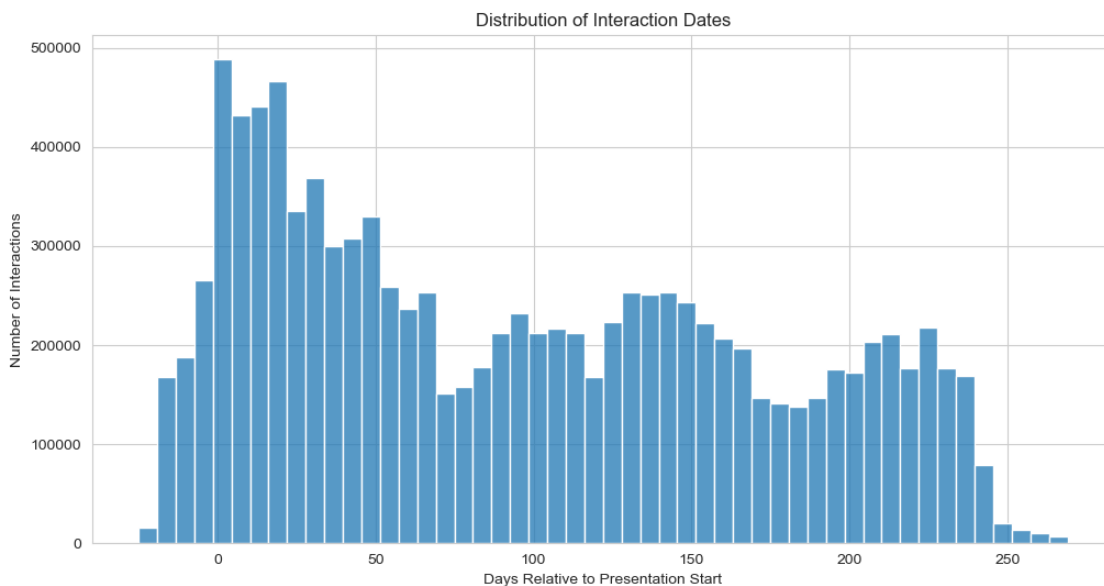
--- Analyzing student_vle ---

Total interaction records: 10655280

Unique students in VLE logs: 26074

Unique VLE items interacted with: 6268

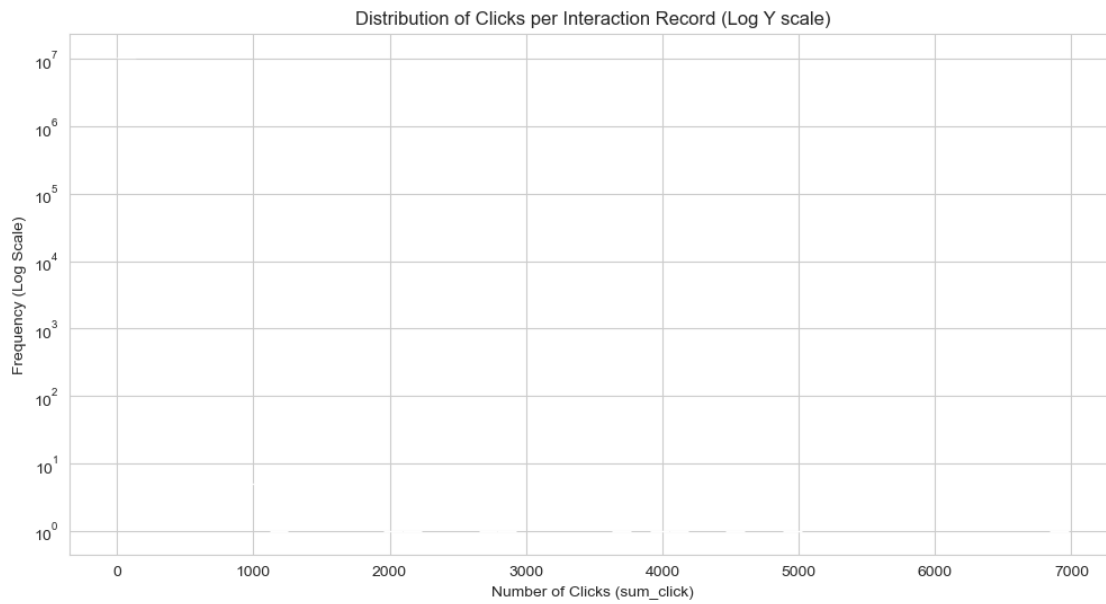
Distribution of Interaction Dates (relative to start):



Distribution of Clicks per Interaction Record:

```
count    1.065528e+07
mean     3.716946e+00
std      8.849047e+00
min      1.000000e+00
25%      1.000000e+00
50%      2.000000e+00
75%      3.000000e+00
90%      8.000000e+00
95%      1.200000e+01
99%      3.400000e+01
max      6.977000e+03
```

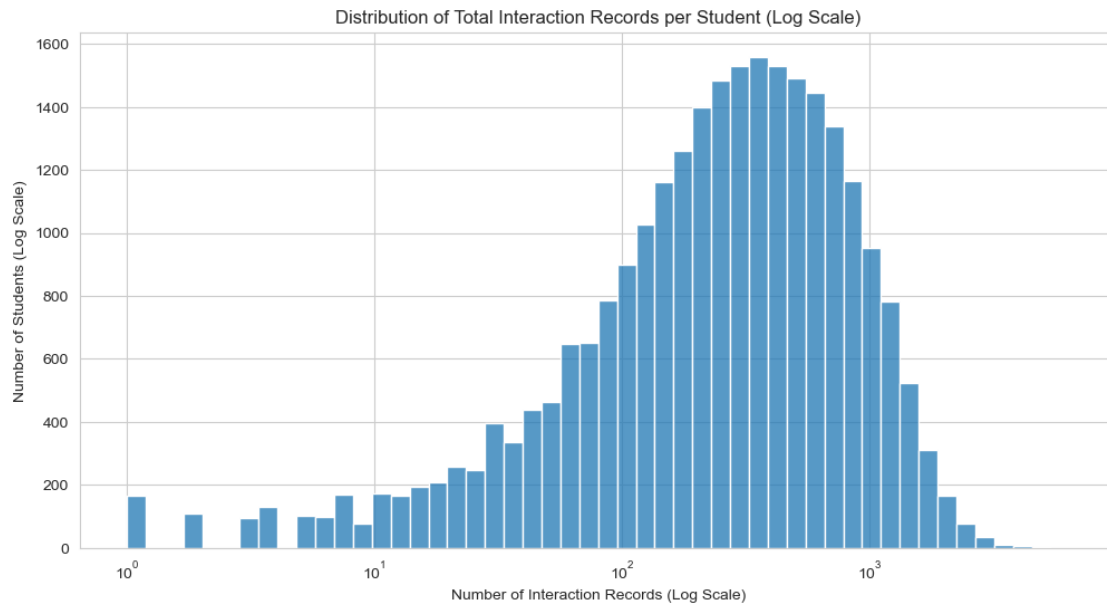
Name: sum_click, dtype: float64



Interactions per Student (Summary Stats):

```
count    26074.000000
mean     408.655365
std      430.608121
min       1.000000
25%      108.000000
50%      270.000000
75%      570.000000
max      6389.000000
```

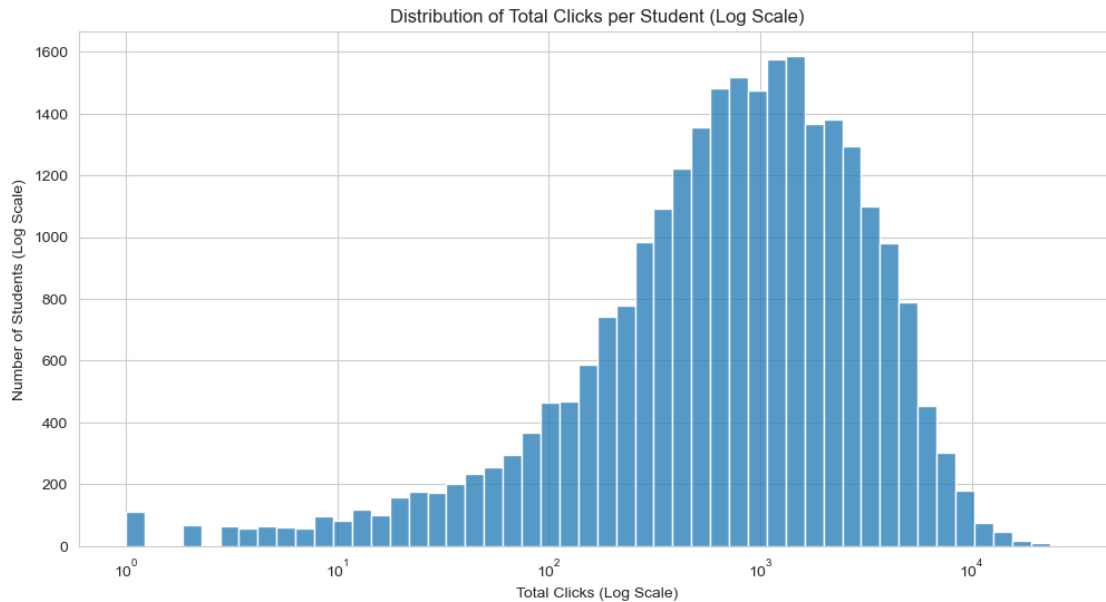
dtype: float64



Total Clicks per Student (Summary Stats):

```
count    26074.000000
mean      1518.949873
std       1935.994635
min        1.000000
25%       298.000000
50%       824.000000
75%      2018.000000
max      28615.000000
```

Name: sum_click, dtype: float64



```
[8]: print("--- Analyzing vle ---")
vle = dataframes['vle']

print(f"Total VLE items defined: {vle.shape[0]}")
print(f"Unique VLE item IDs (id_site): {vle['id_site'].nunique()}") # Should
    ↪ match shape[0]

# Distribution of activity types
print("\nDistribution of VLE Activity Types:")
activity_counts = vle['activity_type'].value_counts()
sns.barplot(y=activity_counts.index, x=activity_counts.values)
plt.title('Distribution of VLE Activity Types')
plt.xlabel('Count')
plt.ylabel('Activity Type')
plt.show()

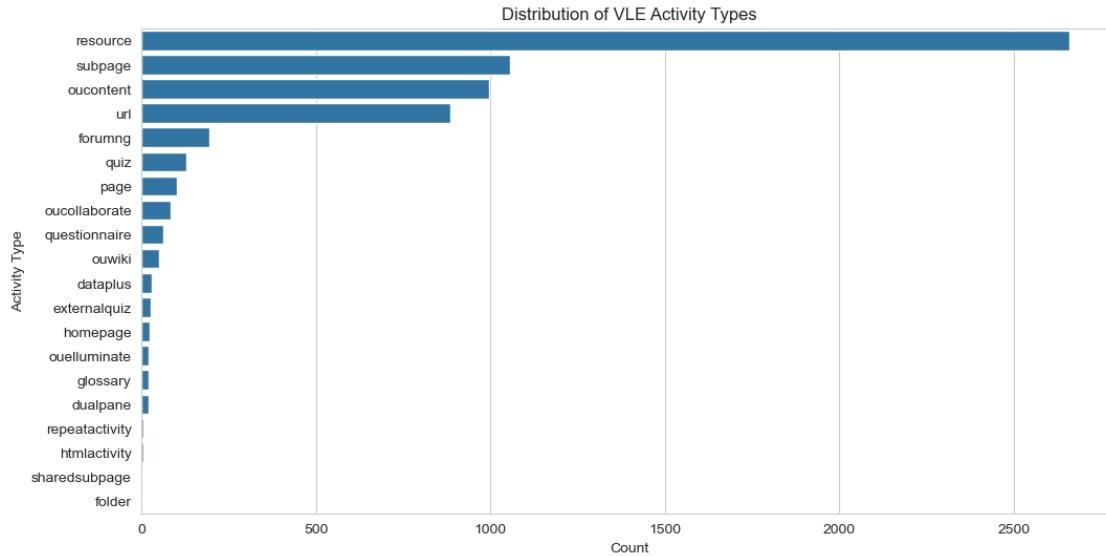
# Weeks - presence of week_from/week_to
print(f"\nPercentage of VLE items with 'week_from': {vle['week_from'].notnull().
    ↪ mean()*100:.2f}%")
print(f"Percentage of VLE items with 'week_to': {vle['week_to'].notnull().
    ↪ mean()*100:.2f}%")
```

--- Analyzing vle ---

Total VLE items defined: 6364

Unique VLE item IDs (id_site): 6364

Distribution of VLE Activity Types:



Percentage of VLE items with 'week_from': 17.61%

Percentage of VLE items with 'week_to': 17.61%

```
[9]: print("--- Analyzing assessments & student_assessment ---")
assessments = dataframes['assessments']
student_assessment = dataframes['student_assessment']

print(f"Total assessments defined: {assessments.shape[0]}")
print(f"Unique assessment IDs: {assessments['id_assessment'].nunique()}") #
    ↳ Should match shape[0]

# Assessment types
print("\nDistribution of Assessment Types:")
sns.countplot(y='assessment_type', data=assessments,
    ↳ order=assessments['assessment_type'].value_counts().index)
plt.title('Distribution of Assessment Types')
plt.show()

# Distribution of assessment weights
sns.histplot(assessments['weight'], bins=20, kde=False)
plt.title('Distribution of Assessment Weights')
plt.xlabel('Weight (%)')
plt.ylabel('Count')
plt.show()

# Explore student scores
print("\nDistribution of Student Scores:")
```

```

print(student_assessment['score'].describe())
sns.histplot(student_assessment['score'].dropna(), bins=20, kde=True)
plt.title('Distribution of Student Assessment Scores')
plt.xlabel('Score')
plt.ylabel('Count')
plt.show()

# Pass/Fail based on score >= 40 (common threshold)
student_assessment['passed'] = student_assessment['score'] >= 40
print("\nPass Rate (based on score >= 40):")
print(student_assessment['passed'].value_counts(normalize=True))

# Submission relative to deadline
assessments_renamed = assessments.rename(columns={'date': 'deadline'})
student_assessment_merged = pd.merge(
    student_assessment,
    assessments_renamed[['id_assessment', 'deadline']],
    on='id_assessment',
    how='left'
)
student_assessment_merged['days_early'] = student_assessment_merged['deadline'] -
    ↪ student_assessment_merged['date_submitted']

print("\nDistribution of Submission Time Relative to Deadline:")
print(student_assessment_merged['days_early'].describe())
sns.histplot(student_assessment_merged['days_early'].dropna(), bins=50)
plt.title('Submission Time Relative to Deadline (Positive = Early)')
plt.xlabel('Days Early')
plt.ylabel('Count')
plt.show()


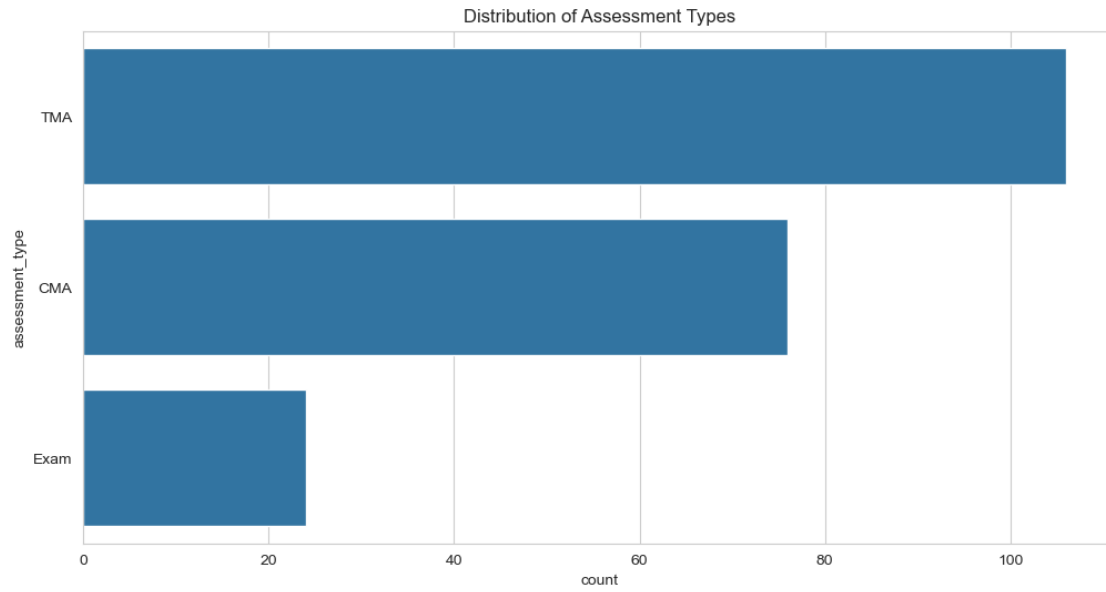
```

--- Analyzing assessments & student_assessment ---

Total assessments defined: 206

Unique assessment IDs: 206

Distribution of Assessment Types:




Distribution of Student Scores:

```
count    173739.000000
mean      75.799573
std       18.798107
min        0.000000
25%       65.000000
50%       80.000000
```

```

75%          90.000000
max          100.000000
Name: score, dtype: float64

```



```

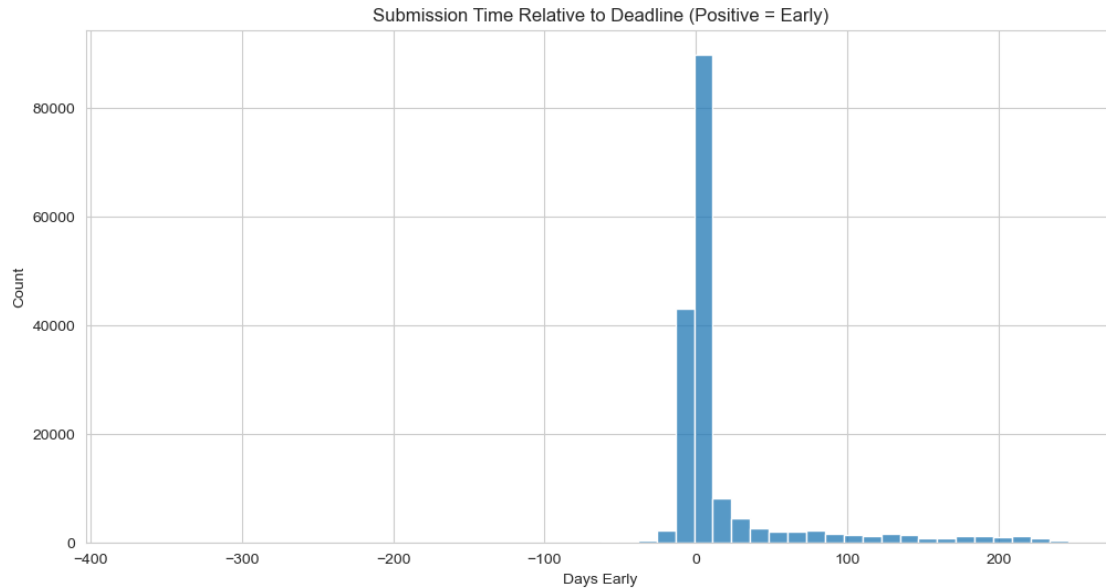
Pass Rate (based on score >= 40):
passed
True      0.955431
False     0.044569
Name: proportion, dtype: float64

```

```

Distribution of Submission Time Relative to Deadline:
count      171047.000000
mean         16.657989
std          45.945880
min         -372.000000
25%          -2.000000
50%           1.000000
75%           6.000000
max          246.000000
Name: days_early, dtype: float64

```



```
[10]: # Example: Average score per highest education level
student_info_subset = student_info[['id_student', 'highest_education']]
student_scores_merged = pd.merge(student_assessment, student_info_subset,
    on='id_student')

avg_score_by_edu = student_scores_merged.groupby('highest_education')['score'].
    mean().sort_values()

print("\nAverage Assessment Score by Highest Education:")
print(avg_score_by_edu)

plt.figure(figsize=(10, 5))
sns.barplot(x=avg_score_by_edu.values, y=avg_score_by_edu.index)
plt.title('Average Assessment Score by Highest Education')
plt.xlabel('Average Score')
plt.ylabel('Highest Education')
plt.show()

# Example: Interactions vs Final Result
student_interactions = student_vle.groupby(['id_student', 'code_module',
    'code_presentation']).size().reset_index(name='total_interactions')
student_clicks = student_vle.groupby(['id_student', 'code_module',
    'code_presentation'])['sum_click'].sum().reset_index(name='total_clicks')

student_activity = pd.merge(student_interactions, student_clicks,
    on=['id_student', 'code_module', 'code_presentation'])
```

```

student_activity_results = pd.merge(student_info[['id_student', 'code_module', 'code_presentation', 'final_result']], student_activity, on=['id_student', 'code_module', 'code_presentation'])

plt.figure(figsize=(12, 6))
sns.boxplot(data=student_activity_results, x='final_result', y='total_clicks', showfliers=False, # Hide outliers for clarity
            order=['Fail', 'Withdrawn', 'Pass', 'Distinction'])
plt.title('Total Clicks vs Final Result (Outliers Hidden)')
plt.ylabel('Total Clicks per Student per Presentation')
plt.yscale('log') # Log scale often useful for clicks
plt.show()

plt.figure(figsize=(12, 6))
sns.boxplot(data=student_activity_results, x='final_result', y='total_interactions', showfliers=False,
            order=['Fail', 'Withdrawn', 'Pass', 'Distinction'])
plt.title('Total Interaction Records vs Final Result (Outliers Hidden)')
plt.ylabel('Total Interaction Records per Student per Presentation')
plt.yscale('log')
plt.show()

```

Average Assessment Score by Highest Education:

highest_education	
No Formal quals	70.601852
Lower Than A Level	73.677280
A Level or Equivalent	75.825197
HE Qualification	77.550154
Post Graduate Qualification	83.489118
Name: score, dtype: float64	

