

## 03\_\_baseline\_\_models

April 21, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import sys
from pathlib import Path
import math # Import math

# Add project root to sys.path
project_root = Path.cwd().parent # Should be RECSYS_FINAL
src_path = project_root / "src"
sys.path.append(str(project_root)) # Add project root for imports like 'src.
    ↪ config'

# Import project modules
from src import config
from src.data import preprocess # For time_based_split
from src.evaluation.evaluator import RecEvaluator # Import the evaluator class
from src.models.popularity import PopularityRecommender # Import the model

# Set display options
pd.set_option('display.max_columns', 100)
pd.set_option('display.max_rows', 100)
sns.set_style("whitegrid")
print("Setup complete. Modules imported.")
print(f"Project Root: {project_root}")
print(f"Processed Data Dir: {config.PROCESSED_DATA_DIR}")
```

```
Loading .env from: /Users/mohit/Desktop/everything/ATLAS/Semester
4/Pinnacle/recsys_final/.env
Database URI configured: Yes
Setup complete. Modules imported.
Project Root: /Users/mohit/Desktop/everything/ATLAS/Semester
4/Pinnacle/recsys_final
Processed Data Dir: /Users/mohit/Desktop/everything/ATLAS/Semester
4/Pinnacle/recsys_final/data/processed
```

```
[2]: # Load the processed parquet files
try:
    interactions_df = pd.read_parquet(config.PROCESSED_DATA_DIR /
    ↪ "interactions_final.parquet")
    users_df = pd.read_parquet(config.PROCESSED_DATA_DIR / "users_final.
    ↪ parquet")
    items_df = pd.read_parquet(config.PROCESSED_DATA_DIR / "items_final.
    ↪ parquet") # Contains presentation_id as column
    print("Processed data loaded successfully.")
    print(f"Interactions shape: {interactions_df.shape}")
    print(f"Users shape: {users_df.shape}")
    print(f"Items shape: {items_df.shape}")

    # Set presentation_id as index for items_df if needed later (evaluator uses
    ↪ it)
    if 'presentation_id' in items_df.columns:
        items_df = items_df.set_index('presentation_id')
        print("Set 'presentation_id' as index for items_df.")

except FileNotFoundError as e:
    print(f"Error loading processed files: {e}")
    print("Please ensure the preprocessing pipeline (run_preprocessing.py) has
    ↪ been run successfully.")
    # Stop execution or handle error
    raise e
except Exception as e:
    print(f"An unexpected error occurred during loading: {e}")
    raise e

# Display heads
print("\nInteractions Head:\n", interactions_df.head())
print("\nUsers Head:\n", users_df.head())
print("\nItems Head:\n", items_df.head())
```

```
Processed data loaded successfully.
Interactions shape: (28466, 7)
Users shape: (25364, 9)
Items shape: (22, 22)
Set 'presentation_id' as index for items_df.
```

```
Interactions Head:
   id_student  presentation_id  total_clicks  interaction_days \
0         6516        AAA_2014J          2791             159
1         8462        DDD_2013J           646              56
2         8462        DDD_2014J           10               1
3        11391        AAA_2013J           934             40
4        23629        BBB_2013B           161             16
```

	first_interaction_date	last_interaction_date	implicit_feedback
0	-23	269	7.934513
1	-6	118	6.472346
2	10	10	2.397895
3	-5	253	6.840547
4	-6	87	5.087596

Users Head:

	id_student	num_of_prev_attempts	studied_credits	gender_mapped	\
0	6516	0	60	0	
1	8462	1	60	0	
2	11391	0	240	0	
3	23629	2	60	1	
4	23698	0	120	1	

	highest_education_mapped	imd_band_mapped	age_band_mapped	\
0	3	9	2	
1	3	4	2	
2	3	10	2	
3	1	3	0	
4	2	6	0	

	disability_mapped	region
0	0	Scotland
1	0	London Region
2	0	East Anglian Region
3	0	East Anglian Region
4	0	East Anglian Region

Items Head:

	module_presentation_length	vle_prop_dataplus	\
presentation_id			
AAA_2013J	268	0.018957	
AAA_2014J	269	0.019802	
BBB_2013J	268	0.000000	
BBB_2014J	262	0.000000	
BBB_2013B	240	0.000000	

	vle_prop_dualpane	vle_prop_externalquiz	vle_prop_folder	\
presentation_id				
AAA_2013J	0.0	0.0	0.0	
AAA_2014J	0.0	0.0	0.0	
BBB_2013J	0.0	0.0	0.0	
BBB_2014J	0.0	0.0	0.0	
BBB_2013B	0.0	0.0	0.0	

	vle_prop_forumng	vle_prop_glossary	vle_prop_homepage	\
--	------------------	-------------------	-------------------	---

presentation_id			
AAA_2013J	0.071090	0.009479	0.004739
AAA_2014J	0.029703	0.009901	0.004950
BBB_2013J	0.059190	0.003115	0.003115
BBB_2014J	0.014493	0.009662	0.004831
BBB_2013B	0.053968	0.003175	0.003175

	vle_prop_htmlactivity	vle_prop_oucollaborate	\
presentation_id			
AAA_2013J	0.0	0.009479	
AAA_2014J	0.0	0.009901	
BBB_2013J	0.0	0.006231	
BBB_2014J	0.0	0.014493	
BBB_2013B	0.0	0.000000	

	vle_prop_oucontent	vle_prop_ouilluminate	vle_prop_ouwiki	\
presentation_id				
AAA_2013J	0.322275	0.000000	0.0	
AAA_2014J	0.336634	0.000000	0.0	
BBB_2013J	0.009346	0.000000	0.0	
BBB_2014J	0.338164	0.000000	0.0	
BBB_2013B	0.003175	0.003175	0.0	

	vle_prop_page	vle_prop_questionnaire	vle_prop_quiz	\
presentation_id				
AAA_2013J	0.0	0.000000	0.000000	
AAA_2014J	0.0	0.000000	0.000000	
BBB_2013J	0.0	0.000000	0.015576	
BBB_2014J	0.0	0.019324	0.019324	
BBB_2013B	0.0	0.000000	0.015873	

	vle_prop_repeatactivity	vle_prop_resource	\
presentation_id			
AAA_2013J	0.0	0.450237	
AAA_2014J	0.0	0.460396	
BBB_2013J	0.0	0.735202	
BBB_2014J	0.0	0.502415	
BBB_2013B	0.0	0.749206	

	vle_prop_sharedsubpage	vle_prop_subpage	vle_prop_url
presentation_id			
AAA_2013J	0.000000	0.028436	0.085308
AAA_2014J	0.000000	0.029703	0.099010
BBB_2013J	0.003115	0.118380	0.046729
BBB_2014J	0.000000	0.048309	0.028986
BBB_2013B	0.003175	0.117460	0.047619

```
[3]: # Cell [3]: Time-Based Split (Using Threshold)

time_col = 'last_interaction_date'
user_col_in_df = 'id_student' # Actual column name in interactions_df
item_col_in_df = 'presentation_id' # Actual column name in interactions_df

# --- Determine Threshold ---
print("--- Determining Time Threshold ---")
print(interactions_df[time_col].describe(percentiles=[.75, .8, .85, .9, .95]))
# Choose threshold based on percentiles (e.g., 80th percentile)
# ***** REPLACE 229 WITH YOUR CHOSEN VALUE *****
TIME_THRESHOLD = 250
print(f"Chosen Time Threshold: {TIME_THRESHOLD}")
print("--- End Threshold Determination ---")

# --- Perform Split ---
if time_col not in interactions_df.columns:
    raise ValueError(f"Time column '{time_col}' not found in interactions data.
↪")
if user_col_in_df not in interactions_df.columns:
    raise ValueError(f"User column '{user_col_in_df}' not found in interactions_
↪data.")
if item_col_in_df not in interactions_df.columns:
    raise ValueError(f"Item column '{item_col_in_df}' not found in interactions_
↪data.")
if not pd.api.types.is_numeric_dtype(interactions_df[time_col]):
    raise TypeError(f"Time column '{time_col}' must be numeric.")

train_df, test_df = preprocess.time_based_split(
    interactions_df=interactions_df,
    user_col=user_col_in_df,
    item_col=item_col_in_df,
    time_col=time_col,
    time_unit_threshold=TIME_THRESHOLD # <<< Use the threshold
    # split_ratio=None # Ensure split_ratio is not used
)

# --- Verify Split ---
print(f"\nTrain shape: {train_df.shape}")
print(f"Test shape: {test_df.shape}")
if not test_df.empty:
    print(f"Min time in Train: {train_df[time_col].min()}, Max time in Train:
↪{train_df[time_col].max()}")
    print(f"Min time in Test: {test_df[time_col].min()}, Max time in Test:
↪{test_df[time_col].max()}")
    # Check user/item overlap
```

```

train_users_final = set(train_df[user_col_in_df].unique())
test_users_final = set(test_df[user_col_in_df].unique())
print(f"Users in Train: {len(train_users_final)}, Users in Test:␣
↪{len(test_users_final)}")
print(f"Users ONLY in Test: {len(test_users_final - train_users_final)}") #␣
↪Should be 0 after filtering in split func

train_items_final = set(train_df[item_col_in_df].unique())
test_items_final = set(test_df[item_col_in_df].unique())
print(f"Items in Train: {len(train_items_final)}, Items in Test:␣
↪{len(test_items_final)}")
print(f"Items ONLY in Test: {len(test_items_final - train_items_final)}") #␣
↪Should be 0 after filtering in split func

else:
    print("Warning: Test DataFrame is empty!")

```

--- Determining Time Threshold ---

```

count    28466.000000
mean      180.275662
std       88.679680
min      -25.000000
50%      228.000000
75%      242.000000
80%      250.000000
85%      257.000000
90%      261.000000
95%      266.000000
max       269.000000

```

Name: last\_interaction\_date, dtype: float64

Chosen Time Threshold: 250

--- End Threshold Determination ---

Performing time-based split...

Original interactions shape: (28466, 7)

Splitting based on time threshold: last\_interaction\_date <= 250

Initial train size: 22892, Initial test size: 5574

Filtered 4836 interactions from test set (users/items not in train).

Final Training set shape: (22892, 7)

Final Test set shape: (738, 7)

Users in Train: 20701, Users in Test: 731

Items in Train: 22, Items in Test: 13

Train shape: (22892, 7)

Test shape: (738, 7)

Min time in Train: -25, Max time in Train: 250

Min time in Test: 251, Max time in Test: 269

Users in Train: 20701, Users in Test: 731

Users ONLY in Test: 0  
Items in Train: 22, Items in Test: 13  
Items ONLY in Test: 0

```
[4]: # Cell [4] - Train Popularity Model

# Initialize and train the Popularity model
# Ensure the item_col matches the column name in train_df and test_df
pop_model = PopularityRecommender(
    user_col='id_student',      # <<< Use the actual user column name
    item_col='presentation_id',  # <<< Use the actual item column name
    score_col='implicit_feedback'
)

# Fit the model using the training data
pop_model.fit(train_df)

# (Optional) Test prediction for a sample user/items
if not test_df.empty:
    sample_user = test_df['id_student'].iloc[0]
    sample_items_all = items_df.index.tolist() # Get all unique item IDs from
    ↪ items_df index
    sample_items_subset = np.random.choice(sample_items_all, min(10,
    ↪ len(sample_items_all)), replace=False).tolist() # Ensure not sampling more
    ↪ than available
    print(f"\nTesting prediction for user {sample_user} on items:
    ↪ {sample_items_subset}")
    scores = pop_model.predict(sample_user, sample_items_subset)
    print("Scores (Popularity):", scores)
else:
    print("\nSkipping prediction test as test_df is empty.")
```

Initialized PopularityRecommender

Fitting PopularityRecommender...

Mapped 20701 users and 22 items.

Fit complete. Calculated popularity for 22 items.

Top 5 most popular items: ['FFF\_2013B', 'CCC\_2014B', 'FFF\_2014B', 'BBB\_2013B', 'BBB\_2014J']

Testing prediction for user 29639 on items: ['GGG\_2013J', 'EEE\_2013J', 'DDD\_2014B', 'GGG\_2014J', 'FFF\_2014B', 'DDD\_2013J', 'EEE\_2014J', 'BBB\_2014B', 'AAA\_2013J', 'EEE\_2014B']

Scores (Popularity): [4234.305017098054, 3786.79398624287, 6704.788396041217, 3105.473475817197, 9406.690532184764, 7253.355677703635, 3896.5541039967225, 7208.319667495403, 1359.433747092062, 4083.0891111100705]

```
[5]: # Cell [5] - Evaluate Popularity Model

# Initialize the evaluator
# Ensure items_df has presentation_id as index before passing
if test_df.empty:
    print("\nCannot evaluate model: Test data is empty.")
elif items_df.index.name != 'presentation_id':
    print("\nError: items_df must have 'presentation_id' set as index for evaluator.")
else:
    evaluator = RecEvaluator(
        train_df=train_df,
        test_df=test_df,
        item_features_df=items_df, # Pass items_df with index set
        user_col='id_student',     # <<< Use the actual user column name
        item_col='presentation_id',# <<< Use the actual item column name
        k=config.TOP_K             # Use K from config
    )

    # Evaluate the popularity model
    # Using n_neg_samples can speed things up significantly for evaluation if needed
    # Set n_neg_samples=100 for faster (approximate) evaluation, or None for full evaluation
    print("\n--- Starting Evaluation of Popularity Model ---")
    pop_results = evaluator.evaluate_model(pop_model, n_neg_samples=100)

    print("\nPopularity Model Evaluation Results:")
    print(pop_results)
```

Evaluator initialized with 22 unique candidate items.  
 Stored 20701 training interactions for filtering.  
 Prepared test data for 731 users.

--- Starting Evaluation of Popularity Model ---

--- Evaluating Model: PopularityRecommender ---

Total test users: 731. Evaluating 731 users known by the model.

Evaluating users: 0%| | 0/731 [00:00<?, ?it/s]

--- Evaluation Results (K=10) ---

Precision@10: 0.0621

Recall@10: 0.6156

NDCG@10: 0.2153

n\_users\_evaluated: 731.0000

n\_users\_skipped: 0.0000



-----  
Popularity Model Evaluation Results:

```
{'Precision@10': 0.06210670314637483, 'Recall@10': 0.615595075239398, 'NDCG@10':  
0.2153109329329855, 'n_users_evaluated': 731}
```

[6]: # Cell [6] - Train ItemCF Model

```
# Import the model
from src.models.item_cf import ItemCFRecommender

# Initialize and train the ItemCF model
itemcf_model = ItemCFRecommender(
    user_col='id_student',      # Use the actual user column name
    item_col='presentation_id', # Use the actual item column name
    score_col='implicit_feedback'
)

# Fit the model using the training data
itemcf_model.fit(train_df)

# (Optional) Test prediction for a sample user/items
if not test_df.empty:
    # Use the same sample user as before or pick a new one
    sample_user_id = test_df['id_student'].iloc[0]
    # Ensure the user exists in the model's mapping
    if sample_user_id in itemcf_model.user_id_to_idx:
        # Get items the user interacted with in train and test for context
        user_train_interactions = train_df[train_df['id_student'] == sample_user_id][
            'presentation_id'].tolist()
        user_test_interactions = test_df[test_df['id_student'] == sample_user_id][
            'presentation_id'].tolist()
        print(f"\n--- ItemCF Prediction Test ---")
        print(f"Sample User ID: {sample_user_id}")
        print(f"User's Training Items: {user_train_interactions}")
        print(f"User's Test Items (Ground Truth): {user_test_interactions}")

        # Predict scores for the test items and a few others
        sample_items_all = items_df.index.tolist()
        items_to_predict = user_test_interactions + np.random.
        choice(sample_items_all, 5, replace=False).tolist()
        items_to_predict = list(set(items_to_predict)) # Ensure unique items

        print(f"Predicting for Items: {items_to_predict}")
        scores = itemcf_model.predict(sample_user_id, items_to_predict)
        print("Predicted Scores:", scores)
        print("--- End Prediction Test ---")
```

```

    else:
        print(f"Sample user {sample_user_id} not found in ItemCF model training_
↳data.")

else:
    print("\nSkipping ItemCF prediction test as test_df is empty.")

```

```

Initialized ItemCFRecommender
Fitting ItemCFRecommender...
Mapped 20701 users and 22 items.
Creating user-item interaction matrix...
Created sparse matrix with shape: (20701, 22) and density: 0.0503
Calculating item-item cosine similarity...
Calculated item similarity matrix shape: (22, 22)
Stored training interactions for prediction filtering.
Fit complete.

```

```

--- ItemCF Prediction Test ---
Sample User ID: 29639
User's Training Items: ['EEE_2014B']
User's Test Items (Ground Truth): ['CCC_2014J']
Predicting for Items: ['CCC_2014J', 'FFF_2014J', 'EEE_2014B', 'AAA_2014J',
'DDD_2014B', 'FFF_2014B']
Predicted Scores: [0.593949099458515, 0.020948052775677548, 0.0, 0.0, 0.0, 0.0]
--- End Prediction Test ---

```

```

[7]: # Cell [7] - Evaluate ItemCF Model

# Evaluate the ItemCF model using the same evaluator instance
if 'evaluator' in locals() and evaluator is not None: # Check if evaluator_
↳exists
    print("\n--- Starting Evaluation of ItemCF Model ---")
    itemcf_results = evaluator.evaluate_model(itemcf_model, n_neg_samples=100)_
↳# Use negative sampling

    print("\nItemCF Model Evaluation Results:")
    print(itemcf_results)

elif test_df.empty:
    print("\nCannot evaluate model: Test data is empty.")
else:
    print("\nError: Evaluator not initialized. Please run Cell [5]_
↳successfully first.")

```

```

--- Starting Evaluation of ItemCF Model ---

```

```
--- Evaluating Model: ItemCFRecommender ---
Total test users: 731. Evaluating 731 users known by the model.
Evaluating users:  0%|          | 0/731 [00:00<?, ?it/s]
```

```
--- Evaluation Results (K=10) ---
Precision@10: 0.0988
Recall@10: 0.9781
NDCG@10: 0.6153
n_users_evaluated: 731.0000
n_users_skipped: 0.0000
-----
```

```
ItemCF Model Evaluation Results:
{'Precision@10': 0.09876880984952119, 'Recall@10': 0.9781121751025992,
'n_users_evaluated': 731}
```

```
[8]: # Cell [10] - Train ItemCF Model

# Import the model
from src.models.item_cf import ItemCFRecommender

print("\n--- Training ItemCF Model ---")

# Initialize and train the ItemCF model
itemcf_model = ItemCFRecommender(
    user_col='id_student',          # Use the actual user column name from
    ↪ interactions_df
    item_col='presentation_id',     # Use the actual item column name from
    ↪ interactions_df
    score_col='implicit_feedback'
)

# Fit the model using the training data
# This might take a moment as it calculates the similarity matrix
itemcf_model.fit(train_df)

# (Optional) Test prediction for a sample user/items
if not test_df.empty:
    # Use the same sample user as before or pick a new one from the test set
    sample_user_id = test_df['id_student'].iloc[0] # Example: first user in
    ↪ test set

    # Ensure the user exists in the model's mapping
    if sample_user_id in itemcf_model.user_id_to_idx:
        # Get items the user interacted with in train and test for context
```

```

        user_train_interactions = train_df[train_df['id_student'] ==
↪sample_user_id]['presentation_id'].tolist()
        user_test_interactions = test_df[test_df['id_student'] ==
↪sample_user_id]['presentation_id'].tolist()
        print(f"\n--- ItemCF Prediction Test ---")
        print(f"Sample User ID: {sample_user_id}")
        print(f"User's Training Items: {user_train_interactions}")
        print(f"User's Test Items (Ground Truth): {user_test_interactions}")

        # Predict scores for the test items and a few others
        sample_items_all = items_df.index.tolist()
        items_to_predict = user_test_interactions + np.random.
↪choice(sample_items_all, 5, replace=False).tolist()
        items_to_predict = list(set(items_to_predict)) # Ensure unique items

        print(f" Predicting for Items: {items_to_predict}")
        scores = itemcf_model.predict(sample_user_id, items_to_predict)
        print(" Predicted Scores:", scores)
        # Display scores alongside item IDs for better readability
        scored_preds = sorted(list(zip(items_to_predict, scores)), key=lambda x:
↪ x[1], reverse=True)
        print(" Predicted Scores (Sorted):", scored_preds)
        print("--- End Prediction Test ---")
    else:
        print(f"\nSample user {sample_user_id} not found in ItemCF model
↪training data (this shouldn't happen if test set was filtered correctly).")
    else:
        print("\nSkipping ItemCF prediction test as test_df is empty.")

print("\n--- Finished Training ItemCF Model ---")

```

```

--- Training ItemCF Model ---
Initialized ItemCFRecommender
Fitting ItemCFRecommender...
Mapped 20701 users and 22 items.
Creating user-item interaction matrix...
Created sparse matrix with shape: (20701, 22) and density: 0.0503
Calculating item-item cosine similarity...
Calculated item similarity matrix shape: (22, 22)
Stored training interactions for prediction filtering.
Fit complete.

--- ItemCF Prediction Test ---
Sample User ID: 29639
User's Training Items: ['EEE_2014B']
User's Test Items (Ground Truth): ['CCC_2014J']

```

```

Predicting for Items: ['CCC_2014J', 'FFF_2013J', 'EEE_2013J', 'DDD_2013B',
'GGG_2014B', 'FFF_2014B']
Predicted Scores: [0.593949099458515, 0.008426447318283632, 0.0645435753807183,
0.016319825324461853, 0.0, 0.0]
Predicted Scores (Sorted): [('CCC_2014J', 0.593949099458515), ('EEE_2013J',
0.0645435753807183), ('DDD_2013B', 0.016319825324461853), ('FFF_2013J',
0.008426447318283632), ('GGG_2014B', 0.0), ('FFF_2014B', 0.0)]
--- End Prediction Test ---

--- Finished Training ItemCF Model ---

```

```

[9]: # Cell [11] - Evaluate ItemCF Model

# Evaluate the ItemCF model using the same evaluator instance
if 'evaluator' in locals() and evaluator is not None: # Check if evaluator
    exists
    print("\n--- Starting Evaluation of ItemCF Model ---")
    itemcf_results = evaluator.evaluate_model(itemcf_model, n_neg_samples=100)
    # Use negative sampling

    print("\nItemCF Model Evaluation Results:")
    print(itemcf_results)

elif test_df.empty:
    print("\nCannot evaluate ItemCF model: Test data is empty.")
else:
    print("\nError: Evaluator not initialized. Please run the cell that
    initializes 'evaluator' successfully first.")

```

```

--- Starting Evaluation of ItemCF Model ---

```

```

--- Evaluating Model: ItemCFRecommender ---

```

```

Total test users: 731. Evaluating 731 users known by the model.

```

```

Evaluating users:  0%|          | 0/731 [00:00<?, ?it/s]

```

```

--- Evaluation Results (K=10) ---

```

```

Precision@10: 0.0988

```

```

Recall@10: 0.9781

```

```

NDCG@10: 0.6153

```

```

n_users_evaluated: 731.0000

```

```

n_users_skipped: 0.0000

```

```

ItemCF Model Evaluation Results:

```

```

{'Precision@10': 0.09876880984952119, 'Recall@10': 0.9781121751025992,
'NDCG@10': 0.6152957703109161, 'n_users_evaluated': 731}

```

```

[10]: # Cell [12] - Train ALS Model

# Import the model
from src.models.matrix_factorization import ImplicitALSWrapper

print("\n--- Training Implicit ALS Model ---")

# Initialize and train the ALS model
# Adjust hyperparameters as needed (these are examples)
als_model = ImplicitALSWrapper(
    user_col='id_student',
    item_col='presentation_id',
    score_col='implicit_feedback',
    factors=50,           # Latent factors
    regularization=0.05,  # Regularization
    iterations=25,        # Iterations
    random_state=config.RANDOM_SEED
)

# Fit the model using the training data
# This will take longer than Popularity or ItemCF
als_model.fit(train_df)

# (Optional) Test prediction for a sample user/items
if not test_df.empty:
    sample_user_id = test_df['id_student'].iloc[0]
    if sample_user_id in als_model.user_id_to_idx:
        user_train_interactions = train_df[train_df['id_student'] == sample_user_id][
            'presentation_id'].tolist()
        user_test_interactions = test_df[test_df['id_student'] == sample_user_id][
            'presentation_id'].tolist()

        print(f"\n--- ALS Prediction Test ---")
        print(f"Sample User ID: {sample_user_id}")
        print(f"User's Training Items: {user_train_interactions}")
        print(f"User's Test Items (Ground Truth): {user_test_interactions}")

        sample_items_all = items_df.index.tolist()
        items_to_predict = user_test_interactions + np.random.
            choice(sample_items_all, 5, replace=False).tolist()
        items_to_predict = list(set(items_to_predict))

        print(f"Predicting for Items: {items_to_predict}")
        scores = als_model.predict(sample_user_id, items_to_predict)
        scored_preds = sorted(list(zip(items_to_predict, scores)), key=lambda x:
            x[1], reverse=True)
        print("Predicted Scores (Sorted):", scored_preds)
        print("--- End Prediction Test ---")

```

```

    else:
        print(f"\nSample user {sample_user_id} not found in ALS model training_
↳data.")
    else:
        print("\nSkipping ALS prediction test as test_df is empty.")

print("\n--- Finished Training Implicit ALS Model ---")

```

```

--- Training Implicit ALS Model ---
Initialized ImplicitALSWrapper
Fitting ImplicitALSWrapper...
  Mapped 20701 users and 22 items.
Creating user-item interaction matrix for Implicit ALS...
  Created sparse matrix (Users x Items) shape: (20701, 22) density: 0.0503
Initializing implicit.als.AlternatingLeastSquares...
Fitting model on User x Item matrix shape: (20701, 22)...

/opt/anaconda3/lib/python3.12/site-packages/implicit/cpu/als.py:95:
RuntimeWarning: OpenBLAS is configured to use 16 threads. It is highly
recommended to disable its internal threadpool by setting the environment
variable 'OPENBLAS_NUM_THREADS=1' or by calling
'threadpoolctl.threadpool_limits(1, "blas")'. Having OpenBLAS use a threadpool
can lead to severe performance issues here.
  check_blas_config()

0%|          | 0/25 [00:00<?, ?it/s]

Model fitting complete.

--- ALS Prediction Test ---
Sample User ID: 29639
User's Training Items: ['EEE_2014B']
User's Test Items (Ground Truth): ['CCC_2014J']
Predicting for Items: ['CCC_2014J', 'DDD_2013B', 'AAA_2013J', 'EEE_2014J',
'CCC_2014B', 'FFF_2014B']
Predicted Scores (Sorted): [('CCC_2014B', 0.0017379806376993656), ('EEE_2014J',
0.0004142490215599537), ('CCC_2014J', 0.00020702210895251483), ('DDD_2013B',
0.00016349671932402998), ('AAA_2013J', -6.1340538195509e-06), ('FFF_2014B',
-9.699559450382367e-05)]
--- End Prediction Test ---

--- Finished Training Implicit ALS Model ---

```

```

[11]: # Cell [13] - Evaluate ALS Model

# Evaluate the ALS model using the same evaluator instance
if 'evaluator' in locals() and evaluator is not None:
    print("\n--- Starting Evaluation of Implicit ALS Model ---")

```

```

    als_results = evaluator.evaluate_model(als_model, n_neg_samples=100) # Use
    ↪negative sampling

    print("\nImplicit ALS Model Evaluation Results:")
    print(als_results)

elif test_df.empty:
    print("\nCannot evaluate ALS model: Test data is empty.")
else:
    print("\nError: Evaluator not initialized. Please run the cell that
    ↪initializes 'evaluator' successfully first.")

```

--- Starting Evaluation of Implicit ALS Model ---

--- Evaluating Model: ImplicitALSWrapper ---

Total test users: 731. Evaluating 731 users known by the model.

Evaluating users: 0%| | 0/731 [00:00<?, ?it/s]

--- Evaluation Results (K=10) ---

Precision@10: 0.0685

Recall@10: 0.6778

NDCG@10: 0.3844

n\_users\_evaluated: 731.0000

n\_users\_skipped: 0.0000

-----

Implicit ALS Model Evaluation Results:

```
{'Precision@10': 0.06853625170998631, 'Recall@10': 0.6778385772913816,
'NDCG@10': 0.3844081787750316, 'n_users_evaluated': 731}
```